

# Customer Churn Analysis

May 22, 2024

## 1 Customer Churn Analysis

This is a dataset of an ecommerce company and we have some customers who are churning (leaving).

### 1.1 Goals

1. Perform exploratory analysis of the provided customer data to share insights of the behavior and characteristics of the customers. Make suggestions to help the company with customer retention
2. Build a predictive model to identify customers who are at risk of leaving the company (churn) based on the provided variables. This can help the company take proactive steps to retain these customers and reduce the rate of churn

### 1.2 Data description

1. CustomerID
2. Churn: Churn Flag
3. Tenure: in months
4. PreferredLoginDevice
5. CityTier
6. WarehouseToHome: Distance in between warehouse to home of customer
7. PreferredPaymentMode
8. Gender
9. HourSpendOnApp
10. NumberOfDeviceRegistered
11. PreferredOrderCat
12. SatisfactionScore
13. MaritalStatus
14. NumberOfAddress
15. OrderAmountHikeFromlastYear: Percentage increases in order from last year
16. CouponUsed: Total number of coupon has been used in last month

17. OrderCount: Total number of orders has been places in last month
18. DaySinceLastOrder
19. CashbackAmount: Average cashback in last month
20. Complain: Complain flag - if the customer ever had a complain

### 1.3 Tips

1. Sharing your thoughts and reasoning as you go will help!
2. Feel free to use any libraries like scikit-learn, use stackoverflow. Don't use ChatGPT and similar.
3. If you are unable to complete any step I can provide help towards the answer. Demonstrating understanding of the solution and result will earn some points!

### 1.4 Import the libraries

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

## 2 1. Data overview

**Question:** Read the sheet named 'E Comm' from file 'E Commerce Dataset.xlsx' saved in current directory into df variable. Print first 5 rows of the dataframe

```
[2]: df = pd.read_excel('E Commerce Dataset.xlsx')
df.head(5)
```

```
[2]:
```

	CustomerID	Churn	Tenure	PreferredLoginDevice	CityTier	WarehouseToHome	\
0	50001	1	4.0	Mobile Phone	3	6.0	
1	50002	1	NaN	Phone	1	8.0	
2	50003	1	NaN	Phone	1	30.0	
3	50004	1	0.0	Phone	3	15.0	
4	50005	1	0.0	Phone	1	12.0	

	PreferredPaymentMode	Gender	HourSpendOnApp	NumberOfDeviceRegistered	\
0	Debit Card	Female	3.0	3	
1	UPI	Male	3.0	4	
2	Debit Card	Male	2.0	4	
3	Debit Card	Male	2.0	4	
4	CC	Male	NaN	3	

	PreferredOrderCat	SatisfactionScore	MaritalStatus	NumberOfAddress	\
0	Laptop & Accessory	2	Single	9	
1	Mobile	3	Single	7	
2	Mobile	3	Single	6	
3	Laptop & Accessory	5	Single	8	

	4	Mobile	5	Single	3
		Complain	OrderAmountHikeFromlastYear	CouponUsed	OrderCount \
0		1	11.0	1.0	1.0
1		1	15.0	0.0	1.0
2		1	14.0	0.0	1.0
3		0	23.0	0.0	1.0
4		0	11.0	1.0	1.0

		DaySinceLastOrder	CashbackAmount
0		5.0	159.93
1		0.0	120.90
2		3.0	120.28
3		3.0	134.07
4		3.0	129.60

```
[3]: df.shape
```

```
[3]: (5630, 20)
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5630 entries, 0 to 5629
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                           5630 non-null   int64
1   Churn                                5630 non-null   int64
2   Tenure                               5366 non-null   float64
3   PreferredLoginDevice                  5630 non-null   object
4   CityTier                              5630 non-null   int64
5   WarehouseToHome                       5379 non-null   float64
6   PreferredPaymentMode                   5630 non-null   object
7   Gender                                5630 non-null   object
8   HourSpendOnApp                         5375 non-null   float64
9   NumberOfDeviceRegistered               5630 non-null   int64
10  PreferedOrderCat                       5630 non-null   object
11  SatisfactionScore                     5630 non-null   int64
12  MaritalStatus                         5630 non-null   object
13  NumberOfAddress                       5630 non-null   int64
14  Complain                              5630 non-null   int64
15  OrderAmountHikeFromlastYear           5365 non-null   float64
16  CouponUsed                            5374 non-null   float64
17  OrderCount                            5372 non-null   float64
18  DaySinceLastOrder                     5323 non-null   float64
19  CashbackAmount                        5630 non-null   float64
```

```
dtypes: float64(8), int64(7), object(5)
memory usage: 879.8+ KB
```

**Question: How many unique values are in each column?**

```
[5]: for columns in df:
      print(columns,df[columns].nunique())
```

```
CustomerID 5630
Churn 2
Tenure 36
PreferredLoginDevice 3
CityTier 3
WarehouseToHome 34
PreferredPaymentMode 7
Gender 2
HourSpendOnApp 6
NumberOfDeviceRegistered 6
PreferedOrderCat 6
SatisfactionScore 5
MaritalStatus 3
NumberOfAddress 15
Complain 2
OrderAmountHikeFromlastYear 16
CouponUsed 17
OrderCount 16
DaySinceLastOrder 22
CashbackAmount 2586
```

**Question: Calculate average churn rate** Note: Churn = 1 means customer has churned

```
[6]: df[df['Churn'] == 1].count()
```

```
[6]: CustomerID          948
      Churn              948
      Tenure            867
      PreferredLoginDevice 948
      CityTier          948
      WarehouseToHome    864
      PreferredPaymentMode 948
      Gender            948
      HourSpendOnApp     890
      NumberOfDeviceRegistered 948
      PreferedOrderCat    948
      SatisfactionScore  948
      MaritalStatus      948
      NumberOfAddress    948
      Complain           948
```

```
OrderAmountHikeFromlastYear    934
CouponUsed                      940
OrderCount                     930
DaySinceLastOrder              894
CashbackAmount                 948
dtype: int64
```

```
[7]: #check if there is any n/a
df['Churn'].isna().nunique()
```

```
[7]: 1
```

```
[8]: #churn rate
df[df['Churn'] == 1].shape[0]/df.shape[0]*100
```

```
[8]: 16.838365896980463
```

**Question: How many missing values / nulls are there in each column?**

```
[9]: for col in df.columns:
      print(col, df[col].isna().sum())
```

```
CustomerID 0
Churn 0
Tenure 264
PreferredLoginDevice 0
CityTier 0
WarehouseToHome 251
PreferredPaymentMode 0
Gender 0
HourSpendOnApp 255
NumberOfDeviceRegistered 0
PreferedOrderCat 0
SatisfactionScore 0
MaritalStatus 0
NumberOfAddress 0
Complain 0
OrderAmountHikeFromlastYear 265
CouponUsed 256
OrderCount 258
DaySinceLastOrder 307
CashbackAmount 0
```

## 3 2. Exploratory Data Analysis

### 3.1 Univariate analysis

Here we will understand select variables

### 3.1.1 1. Numeric variables

**Question: Histogram** Show histograms for all numeric columns. Describe the each variable's distribution briefly to a business stakeholder

```
[10]: numerical_df =  
↳ df[['CustomerID', 'Churn', 'Tenure', 'CityTier', 'WarehouseToHome', 'HourSpendOnApp', 'NumberOfDe',  
↳ 'OrderAmountHikeFromlastYear', 'CouponUsed', 'OrderCount', 'DaySinceLa',  
numerical_df.head(5)
```

```
[10]: CustomerID  Churn  Tenure  CityTier  WarehouseToHome  HourSpendOnApp  \  
0      50001      1    4.0        3           6.0           3.0  
1      50002      1    NaN        1           8.0           3.0  
2      50003      1    NaN        1          30.0           2.0  
3      50004      1    0.0        3          15.0           2.0  
4      50005      1    0.0        1          12.0           NaN  
  
      NumberOfDeviceRegistered  SatisfactionScore  NumberOfAddress  Complain  \  
0                             3                   2                9         1  
1                             4                   3                7         1  
2                             4                   3                6         1  
3                             4                   5                8         0  
4                             3                   5                3         0  
  
      OrderAmountHikeFromlastYear  CouponUsed  OrderCount  DaySinceLastOrder  \  
0                             11.0           1.0           1.0                5.0  
1                             15.0           0.0           1.0                0.0  
2                             14.0           0.0           1.0                3.0  
3                             23.0           0.0           1.0                3.0  
4                             11.0           1.0           1.0                3.0  
  
      CashbackAmount  
0          159.93  
1          120.90  
2          120.28  
3          134.07  
4          129.60
```


```
[11]: categorical_df =  
↳ df[['PreferredLoginDevice', 'PreferredPaymentMode', 'Gender', 'PreferredOrderCat', 'MaritalStatu',  
categorical_df.head(5)
```

```
[11]: PreferredLoginDevice PreferredPaymentMode  Gender  PreferredOrderCat  \  
0      Mobile Phone      Debit Card  Female  Laptop & Accessory  
1          Phone      UPI      Male      Mobile  
2          Phone      Debit Card  Male      Mobile  
3          Phone      Debit Card  Male  Laptop & Accessory  
4          Phone      CC      Male      Mobile
```

```
MaritalStatus
0      Single
1      Single
2      Single
3      Single
4      Single
```

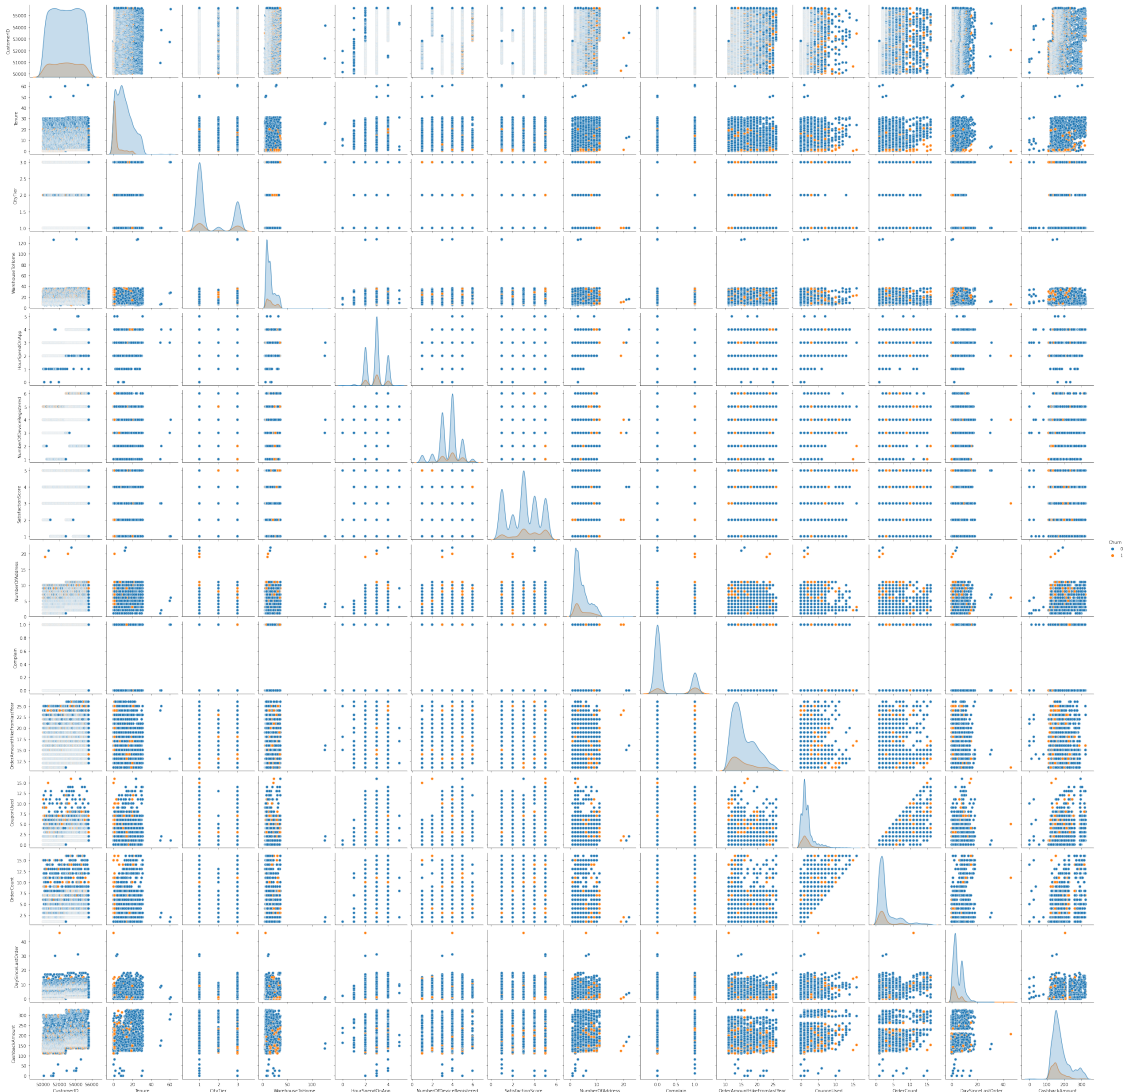
```
[12]: sns.histplot(data=numerical_df, x="CityTier")
```

```
[12]: <AxesSubplot:xlabel='CityTier', ylabel='Count'>
```



```
[13]: sns.pairplot(numerical_df, hue="Churn")
```

```
[13]: <seaborn.axisgrid.PairGrid at 0x7f83c83fb4c0>
```



### 3.1.2 2. Non numeric columns

Cleaning

```
[14]: #As mobile phone and phone are both same so we have merged them
df.loc[df['PreferredLoginDevice'] == 'Phone', 'PreferredLoginDevice'] = 'Mobile Phone'
df.loc[df['PreferredOrderCat'] == 'Mobile', 'PreferredOrderCat'] = 'Mobile Phone'
#as cod is also cash on delivery
#as cc is also credit card so i merged them
df.loc[df['PreferredPaymentMode'] == 'COD', 'PreferredPaymentMode'] = 'Cash on Delivery'
df.loc[df['PreferredPaymentMode'] == 'CC', 'PreferredPaymentMode'] = 'Credit Card'
```



**Question:** Show the unique values of each variable and print the number of occurrences of each value Describe these results briefly

```
[15]: for col in df.columns:
      if df[col].dtype == object:
          print(df[col].unique(), df[col].value_counts())

['Mobile Phone' 'Computer'] Mobile Phone    3996
Computer    1634
Name: PreferredLoginDevice, dtype: int64
['Debit Card' 'UPI' 'Credit Card' 'Cash on Delivery' 'E wallet'] Debit Card
2314
Credit Card    1774
E wallet    614
Cash on Delivery    514
UPI    414
Name: PreferredPaymentMode, dtype: int64
['Female' 'Male'] Male    3384
Female    2246
Name: Gender, dtype: int64
['Laptop & Accessory' 'Mobile Phone' 'Others' 'Fashion' 'Grocery'] Mobile Phone
2080
Laptop & Accessory    2050
Fashion    826
Grocery    410
Others    264
Name: PreferredOrderCat, dtype: int64
['Single' 'Divorced' 'Married'] Married    2986
Single    1796
Divorced    848
Name: MaritalStatus, dtype: int64
```

### 3.2 Analysing the Churn by select variables

Provide business recommendation for each of the below

**Question:** Relation between complains and churn

```
[16]: df['Churn'].corr(df['Complain'])

[16]: 0.25018825469703104

[17]: churn_rate_complain = df[df['Complain'] == 1]['Churn'].mean()
      churn_rate_no_complain = df[df['Complain'] == 0]['Churn'].mean()

[18]: print("Churn rate for Complain =", churn_rate_complain)
      print("Churn rate for No Complain =", churn_rate_no_complain)
```

Churn rate for Complain = 0.3167082294264339  
 Churn rate for No Complain = 0.1092896174863388

### 3.3 Correlation matrix

Visualize the correlation between all variables

**Question: Do we need to do any preprocessing on categorical variables before calculating correlation?**

```
[19]: #One hot coding for categorical dataset
from sklearn.preprocessing import OneHotEncoder
enc = OneHotEncoder()
new_df = enc.fit_transform(categorical_df)
df_ohc = pd.DataFrame(new_df.toarray(), columns=enc.get_feature_names_out(),
dtype=int)
```

```
[20]: df_ohc.head(5)
```

```
[20]: PreferredLoginDevice_Computer PreferredLoginDevice_Mobile Phone \
0 0 1
1 0 0
2 0 0
3 0 0
4 0 0

PreferredLoginDevice_Phone PreferredPaymentMode_CC \
0 0 0
1 1 0
2 1 0
3 1 0
4 1 1

PreferredPaymentMode_COD PreferredPaymentMode_Cash on Delivery \
0 0 0
1 0 0
2 0 0
3 0 0
4 0 0

PreferredPaymentMode_Credit Card PreferredPaymentMode_Debit Card \
0 0 1
1 0 0
2 0 1
3 0 1
4 0 0

PreferredPaymentMode_E wallet PreferredPaymentMode_UPI ... Gender_Male \
```

0	0	0 ...	0
1	0	1 ...	1
2	0	0 ...	1
3	0	0 ...	1
4	0	0 ...	1

	PreferedOrderCat_Fashion	PreferedOrderCat_Grocery \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	PreferedOrderCat_Laptop & Accessory	PreferedOrderCat_Mobile \
0	1	0
1	0	1
2	0	1
3	1	0
4	0	1

	PreferedOrderCat_Mobile Phone	PreferedOrderCat_Others \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	MaritalStatus_Divorced	MaritalStatus_Married	MaritalStatus_Single
0	0	0	1
1	0	0	1
2	0	0	1
3	0	0	1
4	0	0	1

[5 rows x 21 columns]

```
[21]: df_final = pd.concat([numerical_df, df_ohc], axis=1)
df_final.head(5)
```

```
[21]: CustomerID  Churn  Tenure  CityTier  WarehouseToHome  HourSpendOnApp \
0      50001      1    4.0      3          6.0          3.0
1      50002      1    NaN      1          8.0          3.0
2      50003      1    NaN      1         30.0          2.0
3      50004      1    0.0      3         15.0          2.0
4      50005      1    0.0      1         12.0          NaN
```

	NumberOfDeviceRegistered	SatisfactionScore	NumberOfAddress	Complain \
--	--------------------------	-------------------	-----------------	------------

0		3		2		9		1
1		4		3		7		1
2		4		3		6		1
3		4		5		8		0
4		3		5		3		0

	...	Gender_Male	PreferedOrderCat_Fashion	PreferedOrderCat_Grocery	\
0	...	0	0	0	
1	...	1	0	0	
2	...	1	0	0	
3	...	1	0	0	
4	...	1	0	0	

		PreferedOrderCat_Laptop & Accessory	PreferedOrderCat_Mobile	\
0		1	0	
1		0	1	
2		0	1	
3		1	0	
4		0	1	

		PreferedOrderCat_Mobile Phone	PreferedOrderCat_Others	\
0		0	0	
1		0	0	
2		0	0	
3		0	0	
4		0	0	

		MaritalStatus_Divorced	MaritalStatus_Married	MaritalStatus_Single
0		0	0	1
1		0	0	1
2		0	0	1
3		0	0	1
4		0	0	1

[5 rows x 36 columns]

```
[22]: df_final = df_final.dropna()
df_final.head(5)
```

```
[22]: CustomerID  Churn  Tenure  CityTier  WarehouseToHome  HourSpendOnApp  \
0      50001      1    4.0      3          6.0          3.0
3      50004      1    0.0      3          15.0          2.0
5      50006      1    0.0      1          22.0          3.0
11     50012      1   11.0      1           6.0          3.0
12     50013      1    0.0      1          11.0          2.0
```

		NumberOfDeviceRegistered	SatisfactionScore	NumberOfAddress	Complain	\
--	--	--------------------------	-------------------	-----------------	----------	---

0	3	2	9	1
3	4	5	8	0
5	5	5	2	1
11	4	3	10	1
12	3	3	2	1

	Gender_Male	PreferedOrderCat_Fashion	PreferedOrderCat_Grocery	\
0	0	0	0	
3	1	0	0	
5	0	0	0	
11	1	1	0	
12	1	0	0	

	PreferedOrderCat_Laptop & Accessory	PreferedOrderCat_Mobile	\
0	1	0	
3	1	0	
5	0	0	
11	0	0	
12	0	1	

	PreferedOrderCat_Mobile Phone	PreferedOrderCat_Others	\
0	0	0	
3	0	0	
5	1	0	
11	0	0	
12	0	0	

	MaritalStatus_Divorced	MaritalStatus_Married	MaritalStatus_Single
0	0	0	1
3	0	0	1
5	0	0	1
11	0	0	1
12	0	0	1

[5 rows x 36 columns]

**Question: Plot correlation matrix** Discuss a few significant correlations

```
[23]: #CustomerID column as it's not relevant for correlation
df_final = df_final.drop(columns=['CustomerID'])
df_final.head(5)
```

```
[23]:
```

	Churn	Tenure	CityTier	WarehouseToHome	HourSpendOnApp	\
0	1	4.0	3	6.0	3.0	
3	1	0.0	3	15.0	2.0	
5	1	0.0	1	22.0	3.0	
11	1	11.0	1	6.0	3.0	

12	1	0.0	1	11.0	2.0
----	---	-----	---	------	-----

	NumberOfDeviceRegistered	SatisfactionScore	NumberOfAddress	Complain	\
0	3	2	9	1	
3	4	5	8	0	
5	5	5	2	1	
11	4	3	10	1	
12	3	3	2	1	

	OrderAmountHikeFromlastYear	...	Gender_Male	PreferedOrderCat_Fashion	\
0	11.0	...	0	0	
3	23.0	...	1	0	
5	22.0	...	0	0	
11	13.0	...	1	1	
12	13.0	...	1	0	

	PreferedOrderCat_Grocery	PreferedOrderCat_Laptop & Accessory	\
0	0	1	
3	0	1	
5	0	0	
11	0	0	
12	0	0	

	PreferedOrderCat_Mobile	PreferedOrderCat_Mobile Phone	\
0	0	0	
3	0	0	
5	0	1	
11	0	0	
12	1	0	

	PreferedOrderCat_Others	MaritalStatus_Divorced	MaritalStatus_Married	\
0	0	0	0	
3	0	0	0	
5	0	0	0	
11	0	0	0	
12	0	0	0	

	MaritalStatus_Single
0	1
3	1
5	1
11	1
12	1

[5 rows x 35 columns]

```
[24]: corr_matrix = df_final.corr()
      corr_matrix.head(5)
```

```
[24]:
```

	Churn	Tenure	CityTier	WarehouseToHome	\
Churn	1.000000	-0.340013	0.073858	0.087318	
Tenure	-0.340013	1.000000	-0.057414	-0.016353	
CityTier	0.073858	-0.057414	1.000000	0.004457	
WarehouseToHome	0.087318	-0.016353	0.004457	1.000000	
HourSpendOnApp	0.060845	-0.029818	-0.070035	0.052731	

	HourSpendOnApp	NumberOfDeviceRegistered	SatisfactionScore	\
Churn	0.060845	0.149041	0.095759	
Tenure	-0.029818	-0.058752	-0.009972	
CityTier	-0.070035	-0.008616	-0.019494	
WarehouseToHome	0.052731	0.024582	0.000434	
HourSpendOnApp	1.000000	0.293021	0.039879	

	NumberOfAddress	Complain	OrderAmountHikeFromlastYear	...	\
Churn	0.076336	0.238137	0.017193	...	
Tenure	0.196547	-0.035228	0.012197	...	
CityTier	-0.046406	-0.006122	-0.027628	...	
WarehouseToHome	0.003422	0.003829	0.031975	...	
HourSpendOnApp	0.124962	0.020413	0.096827	...	

	Gender_Male	PreferedOrderCat_Fashion	\
Churn	0.033792	0.013086	
Tenure	-0.052123	0.130259	
CityTier	-0.018072	0.083035	
WarehouseToHome	0.035581	-0.009435	
HourSpendOnApp	-0.009952	-0.147830	

	PreferedOrderCat_Grocery	\
Churn	-0.017880	
Tenure	0.091251	
CityTier	-0.001750	
WarehouseToHome	0.003512	
HourSpendOnApp	-0.017382	

	PreferedOrderCat_Laptop & Accessory	PreferedOrderCat_Mobile	\
Churn	-0.184584	0.016675	
Tenure	0.104482	-0.033466	
CityTier	0.245824	-0.097493	
WarehouseToHome	0.043231	-0.015441	
HourSpendOnApp	-0.130819	-0.112896	

	PreferedOrderCat_Mobile Phone	PreferedOrderCat_Others	\
Churn	0.181683	0.010208	

Tenure	-0.211477	0.103702
CityTier	-0.281079	-0.011248
WarehouseToHome	-0.028939	-0.035530
HourSpendOnApp	0.281309	0.023109


	MaritalStatus_Divorced	MaritalStatus_Married \
Churn	-0.013023	-0.159808
Tenure	0.015788	0.085990
CityTier	0.019181	-0.029876
WarehouseToHome	-0.013650	0.038336
HourSpendOnApp	0.018024	0.022257

	MaritalStatus_Single
Churn	0.179481
Tenure	-0.103150
CityTier	0.017370
WarehouseToHome	-0.030496
HourSpendOnApp	-0.037134

[5 rows x 35 columns]

```
[25]: import matplotlib.pyplot as plt
from matplotlib.pyplot import figure
figure(figsize=(10, 8), dpi=150)
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.8)
plt.show()
```





Optional: What is the correlation of each feature with target Sort the correlation in descending order and show the top 5

```
[44]: corr_matrix.sort_values(by = 'Churn', ascending=False).head(5)
```

```
[44]:
```

	Churn	Tenure	CityTier	WarehouseToHome	\
Churn	1.000000	-0.340013	0.073858	0.087318	
Complain	0.238137	-0.035228	-0.006122	0.003829	
PreferredOrderCat_Mobile Phone	0.181683	-0.211477	-0.281079	-0.028939	
MaritalStatus_Single	0.179481	-0.103150	0.017370	-0.030496	
NumberOfDeviceRegistered	0.149041	-0.058752	-0.008616	0.024582	

	HourSpendOnApp	NumberOfDeviceRegistered	\
Churn	0.060845	0.149041	
Complain	0.020413	0.018881	
PreferredOrderCat_Mobile Phone	0.281309	0.188563	

MaritalStatus_Single	-0.037134	0.069345
NumberOfDeviceRegistered	0.293021	1.000000

	SatisfactionScore	NumberOfAddress	Complain \
Churn	0.095759	0.076336	0.238137
Complain	-0.044533	-0.017363	1.000000
PreferedOrderCat_Mobile Phone	0.011297	-0.023102	0.019793
MaritalStatus_Single	-0.025659	0.020125	-0.009180
NumberOfDeviceRegistered	-0.017788	0.066684	0.018881

	OrderAmountHikeFromlastYear	...	Gender_Male \
Churn	0.017193	...	0.033792
Complain	0.003260	...	-0.039522
PreferedOrderCat_Mobile Phone	0.080301	...	0.053910
MaritalStatus_Single	-0.004949	...	-0.027833
NumberOfDeviceRegistered	0.083342	...	-0.026706

	PreferedOrderCat_Fashion \
Churn	0.013086
Complain	-0.000022
PreferedOrderCat_Mobile Phone	-0.253118
MaritalStatus_Single	-0.014253
NumberOfDeviceRegistered	-0.091604

	PreferedOrderCat_Grocery \
Churn	-0.017880
Complain	0.004534
PreferedOrderCat_Mobile Phone	-0.027697
MaritalStatus_Single	-0.013853
NumberOfDeviceRegistered	0.022548

	PreferedOrderCat_Laptop & Accessory \
Churn	-0.184584
Complain	-0.018126
PreferedOrderCat_Mobile Phone	-0.721851
MaritalStatus_Single	-0.060812
NumberOfDeviceRegistered	-0.080608

	PreferedOrderCat_Mobile \
Churn	0.016675
Complain	-0.005327
PreferedOrderCat_Mobile Phone	-0.125238
MaritalStatus_Single	0.037872
NumberOfDeviceRegistered	-0.125343

	PreferedOrderCat_Mobile Phone \
Churn	0.181683

Complain	0.019793
PreferedOrderCat_Mobile Phone	1.000000
MaritalStatus_Single	0.065241
NumberOfDeviceRegistered	0.188563

	PreferedOrderCat_Others \
Churn	0.010208
Complain	0.007866
PreferedOrderCat_Mobile Phone	-0.048049
MaritalStatus_Single	-0.024032
NumberOfDeviceRegistered	0.035364

	MaritalStatus_Divorced	MaritalStatus_Married \
Churn	-0.013023	-0.159808
Complain	0.012777	-0.000364
PreferedOrderCat_Mobile Phone	0.001864	-0.062740
MaritalStatus_Single	-0.288871	-0.737892
NumberOfDeviceRegistered	-0.001099	-0.064517

	MaritalStatus_Single
Churn	0.179481
Complain	-0.009180
PreferedOrderCat_Mobile Phone	0.065241
MaritalStatus_Single	1.000000
NumberOfDeviceRegistered	0.069345

[5 rows x 35 columns]

## 4 3. Modelling

### 4.1 Prepare data

Fill nulls in each column

```
[27]: from sklearn.impute import KNNImputer
      imputer = KNNImputer(n_neighbors=2)
```

```
[28]: for col in df.columns:
      nullCount = df[col].isnull().sum()
      if nullCount > 0:
          df[col]=imputer.fit_transform(df[[col]])
          print("Filled", nullCount, "nulls in", col)
```

Filled 264 nulls in Tenure  
 Filled 251 nulls in WarehouseToHome  
 Filled 255 nulls in HourSpendOnApp  
 Filled 265 nulls in OrderAmountHikeFromlastYear

Filled 256 nulls in CouponUsed  
 Filled 258 nulls in OrderCount  
 Filled 307 nulls in DaySinceLastOrder

**Question: Make the data suitable for model training**

[29]: df\_final.head(5)

```
[29]:
```

	Churn	Tenure	CityTier	WarehouseToHome	HourSpendOnApp	\
0	1	4.0	3	6.0	3.0	
3	1	0.0	3	15.0	2.0	
5	1	0.0	1	22.0	3.0	
11	1	11.0	1	6.0	3.0	
12	1	0.0	1	11.0	2.0	

	NumberOfDeviceRegistered	SatisfactionScore	NumberOfAddress	Complain	\
0	3	2	9	1	
3	4	5	8	0	
5	5	5	2	1	
11	4	3	10	1	
12	3	3	2	1	

	OrderAmountHikeFromlastYear	...	Gender_Male	PreferredOrderCat_Fashion	\
0	11.0	...	0	0	
3	23.0	...	1	0	
5	22.0	...	0	0	
11	13.0	...	1	1	
12	13.0	...	1	0	

	PreferredOrderCat_Grocery	PreferredOrderCat_Laptop & Accessory	\
0	0	1	
3	0	1	
5	0	0	
11	0	0	
12	0	0	

	PreferredOrderCat_Mobile	PreferredOrderCat_Mobile Phone	\
0	0	0	
3	0	0	
5	0	1	
11	0	0	
12	1	0	

	PreferredOrderCat_Others	MaritalStatus_Divorced	MaritalStatus_Married	\
0	0	0	0	
3	0	0	0	
5	0	0	0	
11	0	0	0	

[illegible]

	MaritalStatus_Single
0	1
3	1
5	1
11	1
12	1

```
[5 rows x 35 columns]
```

```
[30]: from sklearn.linear_model import LogisticRegression
x = df_final.drop(['Churn'],axis = 1)
x.head(5)
y = df_final['Churn']
```

[illegible]

```
[32]: #check the shape of X_train and X_test
      x_train.shape, x_test.shape
```

[32]: ((3019, 34), (755, 34))

```
[33]: clf = LogisticRegression(random_state=0).fit(x_train,y_train)
```

```
/Users/jacquie/opt/anaconda3/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

`https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression`

```
n_iter_i = _check_optimize_result(
```

```
[34]: clf.fit(x_train, y_train)
```

```
/Users/jacquie/opt/anaconda3/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[34]: LogisticRegression(random_state=0)
```

```
[35]: #predict result
y_pred = clf.predict(x_test)
```

## 4.2 Model training

**Question: Train one or more models and show their performance on training and test data** Train and show train & test accuracy

```
[36]: from sklearn.metrics import roc_curve, roc_auc_score, classification_report, \
      accuracy_score, confusion_matrix
train_accuracy = accuracy_score(y_pred, y_test)*100
train_accuracy
```

```
[36]: 89.27152317880794
```

```
[37]: y_pred_train = clf.predict(x_train)
y_pred_train
```

```
[37]: array([0, 0, 1, ..., 0, 1, 0])
```

```
[38]: print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, \
      y_pred_train)))
```

Training-set accuracy score: 0.8917

## 5 4. Evaluation

Feel free to use any libraries / stackoverflow

**Question: Show the precision, recall and f1 score for the model with best accuracy**

```
[39]: print(classification_report(y_pred, y_test, digits=6))
```

	precision	recall	f1-score	support
0	0.958333	0.920000	0.938776	675
1	0.495327	0.662500	0.566845	80
accuracy			0.892715	755
macro avg	0.726830	0.791250	0.752810	755

```
weighted avg    0.909273  0.892715  0.899366          755
```

```
print(classification_report(y_train, y_pred_train, digits=6))
```

## 5.1 Confusion matrix

**Question:** Show the confusion matrix Describe the performance of model and steps you could take to improve it

**Bonus:** Plot the confusion matrix

```
[40]: print('Confusion matrix:\n', confusion_matrix(y_pred, y_test))
```

```
Confusion matrix:
[[621  54]
 [ 27  53]]
```

```
[41]: print('Confusion matrix:\n', confusion_matrix(y_pred_train, y_train))
```

```
Confusion matrix:
[[2404  236]
 [  91 288]]
```

## 6 Discussion

```
[42]: #check for overfitting and underfitting
print('Training set score: {:.4f}'.format(clf.score(x_train, y_train)))
```

```
Training set score: 0.8917
```

```
[43]: print('Test set score: {:.4f}'.format(clf.score(x_test, y_test)))
```

```
Test set score: 0.8927
```

The training-set accuracy score is 0.8917 while the test-set accuracy to be 0.8927. These two values are quite comparable. So, there is no question of overfitting.

Next Step, we could perform hyperparameter optimization using gridsearch CV to improve the performance for this particular model

```
[ ]:
```